

Practice Recommendation Systems

In this practice students will implement, in pairs, collaborative filtering algorithm and apply it to a data set scores of movies¹, that we will denominate dataset. Ratings are on a scale of 1 to 5. The data come from $n_u = 943$ users, and a total $n_m = 1682$ film. Students will have a script that will support running, in parts, in each section of the practice, and help them to see if your code gives correct results.

1. Input data

This section describes the script `ex2_cofi.m` that will be used. The first part of this script loads the dataset, which is contained in `ex2_movies.mat`, generating the Y and R matrices in Matlab. The number of rows of the matrix Y is `NUM_MOVIES` and the number of columns is `num_users`. The elements of the matrix Y are the scores $y^{(i,j)}$ (from 1 to 5). The matrix R is a binary indicator matrix, where $R(i, j) = 1$ if the user j scores i film, and $R(i, j) = 0$ otherwise. The purpose of collaborative filtering is to predict, for each user, the scores of movies that has not yet rated, ie $y^{(i,j)}$ values such that $R(i, j) = 0$. After the prediction, we recommend each user movies with the biggest estimated scores.

To better understand the Y matrix, the script `ex2_cofi.m` calculates the average score of the first film (Toy Story) and the screen shows the value obtained.

In this part of the practice matrices X and Θ is also used:

$$X = \begin{bmatrix} \text{---} (x^{(1)})^T \text{---} \\ \text{---} (x^{(2)})^T \text{---} \\ \vdots \\ \text{---} (x^{(n_m)})^T \text{---} \end{bmatrix}, \quad \Theta = \begin{bmatrix} \text{---} (\theta^{(1)})^T \text{---} \\ \text{---} (\theta^{(2)})^T \text{---} \\ \vdots \\ \text{---} (\theta^{(n_u)})^T \text{---} \end{bmatrix}$$

The i th row of X containing the feature vector $x^{(i)}$ of the i th film. The j th row vector Θ contains parameters vector $\theta^{(j)}$ of the user j . Both vectors $x^{(i)}$ and $\theta^{(j)}$ have dimension n . In this exercise we will use $n = 100$, and thus $x^{(i)} \in \mathbb{R}^{100}$ $\theta^{(j)} \in \mathbb{R}^{100}$. Therefore, X is a matrix $n_m \times 100$ and Θ is a matrix $n_u \times 100$.

2 Learning algorithm for collaborative filtering

Now we begin implementing the learning algorithm for collaborative filtering. We start implementing the cost function (without adjustment).

The collaborative filtering algorithm is based on two sets of n -dimensional vectors $x^{(1)}, \dots, x^{(n_m)}$ and $\theta^{(1)}, \dots, \theta^{(n_u)}$, and from which the model predicts, for user j and film i the score $y^{(i,j)} = (\theta^{(j)})^T x^{(i)}$. Given a dataset containing a set with some values $y^{(i,j)}$, the algorithm must learn vector parameters $x^{(1)}, \dots, x^{(n_m)}$ and $\theta^{(1)}, \dots, \theta^{(n_u)}$ that best fit values $y^{(i,j)}$ provided (minimize mean square error).

¹ Movielens 100k Dataset GroupLens Research.

You must complete the code in *cofiCostFunc.m* to calculate the cost function and gradient for collaborative filtering. Parameters are X and θ function, which is what the algorithm seeks to learn. To use the function *fmincg* minimizing the cost function is set to the parameters introduced into a single vector *params*. As you will see, X and θ are extracted from *params*. First we will see how to implement cost and subsequently gradient

2.1 Cost Function collaborative filtering

The cost function of collaborative filtering (without adjustment) is as follows

$$J(x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_u)}) = \frac{1}{2} \sum_{(i,j):r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2$$

Now it is when you change *cofiCostFunc.m* to provide this cost in the variable J . Note that you should only accumulate the cost to the user j and film i if $R(i, j) = 1$. After completing the feature, you can use the *ex2_cofi* script, which runs the cost function. The obtained value should be 22.22.

2.2 Gradient collaborative filtering

Now you must implement the gradient (no regularization). Specifically, you must complete the code *cofiCostFunc.m* to return X_grad and θ_grad variables. Note that X_grad should be an array of the same size as X , and θ_grad must have the same size as θ . The X_grad and θ_grad gradients are given by

$$\frac{\partial J}{\partial x_k^{(i)}} = \sum_{j:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)}) \theta_k^{(j)}$$

$$\frac{\partial J}{\partial \theta_k^{(j)}} = \sum_{i:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)}) x_k^{(i)}$$

Note that the function returns the gradient as a vector. After completing the code gradients through the script *ex2_cofi*, you can check this implementation. If correct, you will see that the numerical and analytical gradient are very similar.

2.3 regularized cost function

The cost function for collaborative filtering with regularization is given by

$$J(x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_u)}) = \frac{1}{2} \sum_{(i,j):r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2 +$$

$$\left(\frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2 \right) + \left(\frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2 \right)$$

regularization must now add the code you've implemented to calculate the cost function J . When you have it ready, you can check if it is correct by *ex2_cofi* script, which runs the cost function. The value that should appear to be close to 31.34.

2.4 regularized Gradient

Now you must implement the gradient of the regularized cost function. Must add the contributions of regularization terms X_grad and Θ_grad as indicated in formulas gradients function regularized cost

$$\frac{\partial J}{\partial x_k^{(i)}} = \sum_{j:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)}) \theta_k^{(j)} + \lambda x_k^{(i)}$$
$$\frac{\partial J}{\partial \theta_k^{(j)}} = \sum_{i:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)}) x_k^{(i)} + \lambda \theta_k^{(j)}$$

After completing the code, you can check if it works correctly by the script `ex2_cofi`.

3 Run the algorithm

After implementing the cost function and gradient, we will train the algorithm and generate recommendations to him. In the next part of the script `ex2_cofi`, you can enter your own ratings of some films, and when the algorithm is executed, will give personalized recommendations. There is a set of films and scored in `ex2_cofi`, but only as an example. Modifies the films scored in `ex2_cofi` according to your own preferences. You can rate the movies you want in the 1682 dataset. The list of films and their number within the dataset are listed in the file `movie_idx.txt`.

Optional 4 Challenge: set the regularization parameter

Enter a code to evaluate different values of the parameter λ . You can use the following values for λ : [0, 0.03, 0.1, 0.3, 1, 3, 10, 30, 100]. For evaluation, divide the array of input data and two submatrices: Y_{train} and Y_{test} where Y_{test} contains approximately 10% of the scores of Y , and Y_{train} the remaining 90%. Y_{train} be used to learn X and Θ with each λ , and Y_{test} to evaluate the performance of the X and Θ learned. As a performance must use the square root of the mean square error:

$$RMSE = \sqrt{\sum_{(i,j):r(i,j)=1} \frac{(y^{(i,j)} - \hat{y}^{(i,j)})^2}{C}}$$

where C is the number of scores in Y_{test} (to evaluate only the films are considered in Y_{test})